


Making a model “speak” its own embedding parameters: setup and the verbalization gap



The same embedding information: a probe can ① decode it (≈ 0.5), yet zero-shot the model can barely ② speak it (≈ 0.01); the difference is the verbalization gap, bridgeable by LoRA fine-tuning.